

Verification of Probing Security for Masked Circuits

A Constructive Approach and Theoretical Study

Kağan Gülsüm

July 21, 2026

Sabancı University — M.Sc. Thesis Defense

Supervisor: *Asst. Prof. Sūha Orhun Mutluergil*

The story

1. Masked circuits defend against side-channel attacks, but masking is easy to get subtly wrong. So, we want to *verify* it, exactly.

The story

1. Masked circuits defend against side-channel attacks, but masking is easy to get subtly wrong. So, we want to *verify* it, exactly.
2. Deciding even one observation is provably hard ($C=P$ -complete), so we cannot count to find an answer.

The story

1. Masked circuits defend against side-channel attacks, but masking is easy to get subtly wrong. So, we want to *verify* it, exactly.
2. Deciding even one observation is provably hard ($C=P$ -complete), so we cannot count to find an answer.
3. Instead, we certify security *constructively*, giving a proof of independence via a *coupling* of the circuit's randomness.

The story

1. Masked circuits defend against side-channel attacks, but masking is easy to get subtly wrong. So, we want to *verify* it, exactly.
2. Deciding even one observation is provably hard ($C=P$ -complete), so we cannot count to find an answer.
3. Instead, we certify security *constructively*, giving a proof of independence via a *coupling* of the circuit's randomness.
4. We present Coppula, a Lean 4 verifier built on this idea, with three components: a rewrite engine, a subset-space traversal, and a probe-set extension.

The story

1. Masked circuits defend against side-channel attacks, but masking is easy to get subtly wrong. So, we want to *verify* it, exactly.
2. Deciding even one observation is provably hard ($C=P$ -complete), so we cannot count to find an answer.
3. Instead, we certify security *constructively*, giving a proof of independence via a *coupling* of the circuit's randomness.
4. We present Coppula, a Lean 4 verifier built on this idea, with three components: a rewrite engine, a subset-space traversal, and a probe-set extension.
5. It verifies a suite of masked gadgets exactly, and the coupling view proves to be helpful.

Motivation & Model


A cipher can be secure on paper and insecure on the bench.

- A real device leaks through power draw, timing, electromagnetic emanation; all of which vary with the *data* it handles.
- Differential power analysis [Kocher et al. '99] correlates many power traces against a guess of an intermediate value, and may recover the key.

So, the insecurity lies not in the mathematics behind the cipher, but in the way the hardware computes it.

Masking — and our running example

Boolean masking [Chari et al. '99] splits each secret into shares, $x = x_0 \oplus x_1$, with x_0 uniformly random, and computation happens on shares, never directly on secrets. Then, the cost of attacking grows exponentially in the number of shares.



DOM-AND [Groß et al. '16]: a masked AND gate, with two shares per input, and a fresh mask r .

- computes a sharing $s_0 \oplus s_1 = x \odot y$,
- the cross terms p_{01}, p_{10} mix shares of *both* secrets — the delicate part,
- r re-masks them before they meet an output.


This gadget accompanies us for the rest of the talk.

The probing model

Reasoning about noisy traces directly is unwieldy. Luckily we have an abstraction [Ishai–Sahai–Wagner '03]:

d -probing security

An adversary picks any d wires and is handed their *exact* values. The circuit is d -probing secure when no such choice distinguishes one secret assignment from another.




The probing model

Reasoning about noisy traces directly is unwieldy. Luckily we have an abstraction [Ishai–Sahai–Wagner '03]:

d -probing security

An adversary picks any d wires and is handed their *exact* values. The circuit is d -probing secure when no such choice distinguishes one secret assignment from another.




- Now, security becomes a *combinatorial* property of the circuit, algorithmically decidable in principle.

The probing model

Reasoning about noisy traces directly is unwieldy. Luckily we have an abstraction [Ishai–Sahai–Wagner '03]:

d -probing security

An adversary picks any d wires and is handed their *exact* values. The circuit is d -probing secure when no such choice distinguishes one secret assignment from another.



- Now, security becomes a *combinatorial* property of the circuit, algorithmically decidable in principle.
- And thankfully, the abstraction is faithful: probing security implies noisy-leakage security [Duc–Dziembowski–Faust '14]. We take d -probing security as the target and do not revisit this bridge.

The problem

Setting. We work with arithmetic circuits over $\mathbb{F}_2$, two-input $\oplus$ and $\odot$ gates, where inputs are partitioned into secrets $\mathcal{S}$, randoms $\mathcal{R}$, publics $\mathcal{P}$. We have value-based leakage, so, no glitches, and also no composition.

Exact verification of d -probing security

One should decide, for every one of the $\sum_{i \leq d} \binom{|\mathcal{W}|}{i}$ observation sets, whether it is independent of the secrets.

- Secure $\Rightarrow$ every observation set accounted for,
- Insecure $\Rightarrow$ a specific leaking observation as witness.

The problem

Setting. We work with arithmetic circuits over $\mathbb{F}_2$, two-input $\oplus$ and $\odot$ gates, where inputs are partitioned into secrets $\mathcal{S}$, randoms $\mathcal{R}$, publics $\mathcal{P}$. We have value-based leakage, so, no glitches, and also no composition.

Exact verification of d -probing security

One should decide, for every one of the $\sum_{i \leq d} \binom{|\mathcal{W}|}{i}$ observation sets, whether it is independent of the secrets.

- Secure $\Rightarrow$ every observation set accounted for,
 - Insecure $\Rightarrow$ a specific leaking observation as witness.
-
- Two difficulties arise: a *single* probe already asks a question about a distribution over all of $\{0, 1\}^{\mathcal{R}}$,

The problem

Setting. We work with arithmetic circuits over $\mathbb{F}_2$, two-input $\oplus$ and $\odot$ gates, where inputs are partitioned into secrets $\mathcal{S}$, randoms $\mathcal{R}$, publics $\mathcal{P}$. We have value-based leakage, so, no glitches, and also no composition.

Exact verification of d -probing security

One should decide, for every one of the $\sum_{i \leq d} \binom{|\mathcal{W}|}{i}$ observation sets, whether it is independent of the secrets.

- Secure $\Rightarrow$ every observation set accounted for,
- Insecure $\Rightarrow$ a specific leaking observation as witness.
- Two difficulties arise: a *single* probe already asks a question about a distribution over all of $\{0, 1\}^{\mathcal{R}}$,
- and there are combinatorially many probes to settle.

How Hard Is One Probe?

Warm-up: three views of secret-independence

Fix a probe O (a tuple w of wires). Write $\mathcal{L}_a(O)$ for its distribution over a uniform tape $\rho \in \{0,1\}^{\mathcal{R}}$, and $N_a(o) = |\{\rho : O(\rho; a) = o\}|$ for the tape counts.

1. **Distributional.** $\mathcal{L}_a(O) = \mathcal{L}_b(O)$ for all secrets a, b . *the definition*

Warm-up: three views of secret-independence

Fix a probe O (a tuple w of wires). Write $\mathcal{L}_a(O)$ for its distribution over a uniform tape $\rho \in \{0,1\}^{\mathcal{R}}$, and $N_a(o) = |\{\rho : O(\rho; a) = o\}|$ for the tape counts.

1. **Distributional.** $\mathcal{L}_a(O) = \mathcal{L}_b(O)$ for all secrets a, b . *the definition*
2. **Counting.** $N_a(o) = N_b(o)$ for every observation o , i.e. two histograms of integers coincide. *→ this section*

Warm-up: three views of secret-independence

Fix a probe O (a tuple w of wires). Write $\mathcal{L}_a(O)$ for its distribution over a uniform tape $\rho \in \{0,1\}^{\mathcal{R}}$, and $N_a(o) = |\{\rho : O(\rho; a) = o\}|$ for the tape counts.

1. **Distributional.** $\mathcal{L}_a(O) = \mathcal{L}_b(O)$ for all secrets a, b . *the definition*
2. **Counting.** $N_a(o) = N_b(o)$ for every observation o , i.e. two histograms of integers coincide. *→ this section*
3. **Coupling.** there exists a bijection φ of the tapes such that $O(\rho; a) = O(\varphi(\rho); b)$.
→ next section

Warm-up: three views of secret-independence

Fix a probe O (a tuple w of wires). Write $\mathcal{L}_a(O)$ for its distribution over a uniform tape $\rho \in \{0, 1\}^{\mathcal{R}}$, and $N_a(o) = |\{\rho : O(\rho; a) = o\}|$ for the tape counts.

1. **Distributional.** $\mathcal{L}_a(O) = \mathcal{L}_b(O)$ for all secrets a, b . *the definition*
2. **Counting.** $N_a(o) = N_b(o)$ for every observation o , i.e. two histograms of integers coincide. *→ this section*
3. **Coupling.** there exists a bijection φ of the tapes such that $O(\rho; a) = O(\varphi(\rho); b)$. *→ next section*

All three are equivalent — but they are *not* equally usable. The counting view tells us what a verification procedure should not attempt directly; the coupling view tells us what it might be able to accomplish.

Counting is intractable

Theorem (thesis, Ch. 4)

Computing one tape count $N_a(o)$ is #P-complete [Valiant '79]. And deciding $\mathcal{L}_a(O) = \mathcal{L}_b(O)$ is C=P-complete [Wagner '86].

Hardness, briefly: for formulas φ_0, φ_1 over randoms $x_1, \dots, x_m$ and a single secret bit s , the wire

$$w = \varphi_0 \oplus (s \odot (\varphi_0 \oplus \varphi_1))$$

reads φ_0 when $s = 0$ and φ_1 when $s = 1$. Then, the one-bit probe $\{w\}$ is secret-independent iff $\#\varphi_0 = \#\varphi_1$. That is exactly the canonical C=P-complete question: EXACTSAT.

Theorem (thesis, Ch. 4)

Computing one tape count $N_a(o)$ is $\#P$ -complete [Valiant '79]. And deciding $\mathcal{L}_a(O) = \mathcal{L}_b(O)$ is $C=P$ -complete [Wagner '86].

Hardness, briefly: for formulas φ_0, φ_1 over randoms $x_1, \dots, x_m$ and a single secret bit s , the wire

$$w = \varphi_0 \oplus (s \odot (\varphi_0 \oplus \varphi_1))$$

reads φ_0 when $s = 0$ and φ_1 when $s = 1$. Then, the one-bit probe $\{w\}$ is secret-independent iff $\#\varphi_0 = \#\varphi_1$. That is exactly the canonical $C=P$ -complete question: EXACTSAT.

Observe that one secret bit and one probed wire already suffice; no feature of masking is what makes the problem hard.

What the placement tells

- Observe that $C=P \subseteq P^{\#P}$, so the decision never costs more than the counting. Another observation is the following: a formula is unsatisfiable exactly when its count is zero, i.e. $\text{coNP} \subseteq C=P$. Hence, no polynomial exact decision exists unless $P = \text{NP}$.

Consequence

We do not compute or compare the counts. We certify security *constructively*: by exhibiting, for each probe, a measure-preserving relabelling of the randomness under which the secrets visibly cancel. A witness proves independence without explicitly counting — at the price of completeness!

What the placement tells

- Observe that $C=P \subseteq P^{\#P}$, so the decision never costs more than the counting. Another observation is the following: a formula is unsatisfiable exactly when its count is zero, i.e. $\text{coNP} \subseteq C=P$. Hence, no polynomial exact decision exists unless $P = \text{NP}$.
- Both classes sit *above the polynomial hierarchy* (due to Toda). A single probe is already there, and full verification universally quantifies over exponentially many such instances.


Consequence

We do not compute or compare the counts. We certify security *constructively*: by exhibiting, for each probe, a measure-preserving relabelling of the randomness under which the secrets visibly cancel. A witness proves independence without explicitly counting — at the price of completeness!

Security Is a Coupling

One substitution, on the running example

Probe the masked cross term: $c_{01} = x_0 \odot y_1 \oplus r$.




Apply the substitution $\sigma = (r \mapsto x_0 \odot y_1 \oplus r)$:

$$c_{01} \sigma = x_0 \odot y_1 \oplus (x_0 \odot y_1 \oplus r) = r.$$

One substitution, on the running example

Probe the masked cross term: $c_{01} = x_0 \odot y_1 \oplus r$.




Apply the substitution $\sigma = (r \mapsto x_0 \odot y_1 \oplus r)$:

$$c_{01} \sigma = x_0 \odot y_1 \oplus (x_0 \odot y_1 \oplus r) = r.$$

- The probe is now *syntactically secret-free*: it is uniform, whatever the secrets are.

One substitution, on the running example

Probe the masked cross term: $c_{01} = x_0 \odot y_1 \oplus r$.




Apply the substitution $\sigma = (r \mapsto x_0 \odot y_1 \oplus r)$:

$$c_{01} \sigma = x_0 \odot y_1 \oplus (x_0 \odot y_1 \oplus r) = r.$$

- The probe is now *syntactically secret-free*: it is uniform, whatever the secrets are.
- And σ is not just a change of variables. It induces a *bijection* $\hat{\sigma}$ of the tape space $\{0, 1\}^{\mathcal{R}}$ by translating one coordinate — a *shear*.

One substitution, on the running example

Probe the masked cross term: $c_{01} = x_0 \odot y_1 \oplus r$.



Apply the substitution $\sigma = (r \mapsto x_0 \odot y_1 \oplus r)$:

$$c_{01} \sigma = x_0 \odot y_1 \oplus (x_0 \odot y_1 \oplus r) = r.$$

- The probe is now *syntactically secret-free*: it is uniform, whatever the secrets are.
- And σ is not just a change of variables. It induces a *bijection* $\hat{\sigma}$ of the tape space $\{0, 1\}^{\mathcal{R}}$ by translating one coordinate — a *shear*.

Key identity: $(w\sigma)(\rho) = w(\hat{\sigma}(\rho))$ — the syntactic step and the semantic relabelling agree.

The classical idea [Barthe et al. '17 for probabilistic programs] to prove that two distributions are equal without explicitly computing them is by *linking their sources*. Our two laws are pushforwards of the *same* uniform tape. So we should cleverly link those tapes.

Definition (coupling function)

A *coupling function* from a to b for a probe O is a bijection $\varphi : \{0, 1\}^{\mathcal{R}} \rightarrow \{0, 1\}^{\mathcal{R}}$ with

$$O(\rho; a) = O(\varphi(\rho); b) \quad \text{for every tape } \rho.$$

The classical idea [Barthe et al. '17 for probabilistic programs] to prove that two distributions are equal without explicitly computing them is by *linking their sources*. Our two laws are pushforwards of the *same* uniform tape. So we should cleverly link those tapes.

Definition (coupling function)

A *coupling function* from a to b for a probe O is a bijection $\varphi : \{0, 1\}^{\mathcal{R}} \rightarrow \{0, 1\}^{\mathcal{R}}$ with

$$O(\rho; a) = O(\varphi(\rho); b) \quad \text{for every tape } \rho.$$

- A bijection of a finite set preserves the uniform law trivially. The identity says that a run under a is reproduced, observation for observation, under b . This is a *diagonal* coupling of the two laws. It proves that the distributions on the observation set are the same for those two secrets.

Security *is* a coupling

Theorem

A probe O is secret-independent *iff* for every ordered pair a, b there is a coupling function from a to b for O .

Proof sketch.

Theorem

A probe O is secret-independent *iff* for every ordered pair a, b there is a coupling function from a to b for O .

Proof sketch.

- ($\Leftarrow$) Reindex by $\rho' = \varphi(\rho)$:

$$N_a(o) = |\{\rho : O(\rho; a) = o\}| = |\{\rho : O(\varphi(\rho); b) = o\}| = N_b(o).$$

Security *is* a coupling

Theorem

A probe O is secret-independent *iff* for every ordered pair a, b there is a coupling function from a to b for O .

Proof sketch.

- ($\Leftarrow$) Reindex by $\rho' = \varphi(\rho)$:

$$N_a(o) = |\{\rho : O(\rho; a) = o\}| = |\{\rho : O(\varphi(\rho); b) = o\}| = N_b(o).$$

- ($\Rightarrow$) Sort the tapes by outcome: $F_a(o) = \{\rho : O(\rho; a) = o\}$. As o varies these *fibres* partition $\{0, 1\}^{\mathcal{R}}$, and equal laws mean matching fibres have equal size, $|F_a(o)| = |F_b(o)|$. So pick, for each o , any bijection $F_a(o) \rightarrow F_b(o)$; their union is a bijection φ of the whole tape space, and it sends every ρ to a tape with the *same* observation, giving us a coupling function. $\square$

Security *is* a coupling

Theorem

A probe O is secret-independent *iff* for every ordered pair a, b there is a coupling function from a to b for O .

Proof sketch.

- ($\Leftarrow$) Reindex by $\rho' = \varphi(\rho)$:

$$N_a(o) = |\{\rho : O(\rho; a) = o\}| = |\{\rho : O(\varphi(\rho); b) = o\}| = N_b(o).$$

- ($\Rightarrow$) Sort the tapes by outcome: $F_a(o) = \{\rho : O(\rho; a) = o\}$. As o varies these *fibres* partition $\{0, 1\}^{\mathcal{R}}$, and equal laws mean matching fibres have equal size, $|F_a(o)| = |F_b(o)|$. So pick, for each o , any bijection $F_a(o) \rightarrow F_b(o)$; their union is a bijection φ of the whole tape space, and it sends every ρ to a tape with the *same* observation, giving us a coupling function. $\square$

This is an existence statement and not a procedure

The ($\Rightarrow$) direction matches fibres *block by block*, and for that it needs the counts $N_a(o)$, which are #P-hard to compute. So, the theorem guarantees a coupling, it does not find one.

From substitutions to couplings — the constructive route

- Each substitution induces a shear, a bijection of the tapes. A *sequence* of substitutions composes to a bijection φ_a — one per secret assignment, since the contexts e_i may mention secrets.

From substitutions to couplings — the constructive route

- Each substitution induces a shear, a bijection of the tapes. A *sequence* of substitutions composes to a bijection φ_a — one per secret assignment, since the contexts e_i may mention secrets.
- If the sequence drives the probe to a secret-free form $\overline{O}$, then *both* runs agree like the following:

$$O(\varphi_a(\rho); a) = \overline{O}(\rho) = O(\varphi_b(\rho); b).$$

From substitutions to couplings — the constructive route

- Each substitution induces a shear, a bijection of the tapes. A *sequence* of substitutions composes to a bijection φ_a — one per secret assignment, since the contexts e_i may mention secrets.
- If the sequence drives the probe to a secret-free form $\overline{O}$, then *both* runs agree like the following:

$$O(\varphi_a(\rho); a) = \overline{O}(\rho) = O(\varphi_b(\rho); b).$$

- Chain through the common form as $\varphi_b \circ \varphi_a^{-1}$, and one gets a coupling function from a to b — for every pair of secrets at once, from one derivation.

From substitutions to couplings — the constructive route

- Each substitution induces a shear, a bijection of the tapes. A *sequence* of substitutions composes to a bijection φ_a — one per secret assignment, since the contexts e_i may mention secrets.
- If the sequence drives the probe to a secret-free form $\overline{O}$, then *both* runs agree like the following:


$$O(\varphi_a(\rho); a) = \overline{O}(\rho) = O(\varphi_b(\rho); b).$$

- Chain through the common form as $\varphi_b \circ \varphi_a^{-1}$, and one gets a coupling function from a to b — for *every* pair of secrets at once, from one derivation.
- So a successful rewrite derivation is not a simple Secure verdict. It is an explicit, checkable, reusable *certificate*.


This is the design principle of Coppula: replace intractable counting by the search for a coupling, building it one substitution at a time.

Coppula: Three Components

Pipeline overview




Pipeline overview




- I decides a *single* probe, by construction.

Pipeline overview



- I decides a *single* probe, by construction.
- II sweeps every d -subset without explicitly enumerating them.

Pipeline overview



- I decides a *single* probe, by construction.
- II sweeps *every* d -subset without explicitly enumerating them.
- III amplifies one certificate into a whole *safe set*, shrinking the space II must sweep.

Component I: the rewrite engine

Decides one probe w by substituting randoms away, one at a time, driving the observation toward a secret-free form.

- **Tiered.** A cheap single-occurrence rule first (r occurs once across w : substitute, done); with a complete substitution loop behind it.

Component I: the rewrite engine

Decides one probe w by substituting randoms away, one at a time, driving the observation toward a secret-free form.

- **Tiered.** A cheap single-occurrence rule first (r occurs once across w : substitute, done); with a complete substitution loop behind it.
- **Sound.** Every step is a shear, so any Secure verdict carries a genuine coupling (Theorem, Ch. 5).

Component I: the rewrite engine

Decides one probe w by substituting randoms away, one at a time, driving the observation toward a secret-free form.

- **Tiered.** A cheap single-occurrence rule first (r occurs once across w : substitute, done); with a complete substitution loop behind it.
- **Sound.** Every step is a shear, so any Secure verdict carries a genuine coupling (Theorem, Ch. 5).
- **Incomplete, knowingly.** A probe can be independent for reasons no relabelling of this kind captures — this is the price of the hardness result.

Component I: the rewrite engine

Decides one probe w by substituting randoms away, one at a time, driving the observation toward a secret-free form.

- **Tiered.** A cheap single-occurrence rule first (r occurs once across w : substitute, done); with a complete substitution loop behind it.
- **Sound.** Every step is a shear, so any Secure verdict carries a genuine coupling (Theorem, Ch. 5).
- **Incomplete, knowingly.** A probe can be independent for reasons no relabelling of this kind captures — this is the price of the hardness result.

On the example: $c_{01} \xrightarrow{r \mapsto a_0 \odot b_1 \oplus r} r$ is a single step discharge, which is Secure, and records a coupling.

Component I: why factoring

Some masks are buried inside *products*, where substitution cannot reach them.

- Substituting $r \mapsto e \oplus r$ needs r to occur as a *summand*; in $u \odot r \oplus u \odot v$ every occurrence of r is trapped inside a product.

Component I: why factoring

Some masks are buried inside *products*, where substitution cannot reach them.

- Substituting $r \mapsto e \oplus r$ needs r to occur as a *summand*; in $u \odot r \oplus u \odot v$ every occurrence of r is trapped inside a product.
- **Multiplicative factoring** pulls the common factor out,

$$u \odot r \oplus u \odot v \longmapsto u \odot (r \oplus v),$$

and the hidden random surfaces as a summand again, ready to substitute.

Component I: why factoring

Some masks are buried inside *products*, where substitution cannot reach them.

- Substituting $r \mapsto e \oplus r$ needs r to occur as a *summand*; in $u \odot r \oplus u \odot v$ every occurrence of r is trapped inside a product.
- **Multiplicative factoring** pulls the common factor out,

$$u \odot r \oplus u \odot v \longmapsto u \odot (r \oplus v),$$

and the hidden random surfaces as a summand again, ready to substitute.

- Factoring is applied lazily, on demand, between substitution attempts; the engine interleaves the two until the probe is discharged or no rule applies.

Component II: covering all probes

Certifying each d -subset in turn is infeasible. Instead the space lives as a *worklist* of factors $\langle (c_1, W_1), \dots, (c_m, W_m) \rangle$, “draw c_j wires from each W_j ”, and we split recursively [in the manner of maskVerif, Barthe et al. '19]:

- Initially, try to discharge a whole batch at once: if the *union* of its wires is jointly secret-free, every probe inside is also settled — prune the branch.

Theorem (exhaustiveness, Ch. 5)

The recursion visits every d -subset exactly once. The proof reads Vandermonde's identity $\binom{m+n}{d} = \sum_k \binom{m}{k} \binom{n}{d-k}$ at the level of *sets*: the cells genuinely partition the probes, they do not merely count them.

Component II: covering all probes

Certifying each d -subset in turn is infeasible. Instead the space lives as a *worklist* of factors $\langle (c_1, W_1), \dots, (c_m, W_m) \rangle$, “draw c_j wires from each W_j ”, and we split recursively [in the manner of maskVerif, Barthe et al. '19]:

- Initially, try to discharge a whole batch at once: if the *union* of its wires is jointly secret-free, every probe inside is also settled — prune the branch.
- On failure, certify one representative *probe* w (one probe of the target size, c_j wires from each factor), grow it into a safe set (Component III), and split every factor along safe / unsafe regions. Recurse on the resulting cells; the all-safe cell needs no visit.

Theorem (exhaustiveness, Ch. 5)

The recursion visits every d -subset exactly once. The proof reads Vandermonde's identity $\binom{m+n}{d} = \sum_k \binom{m}{k} \binom{n}{d-k}$ at the level of *sets*: the cells genuinely partition the probes, they do not merely count them.

Component II: covering all probes

Certifying each d -subset in turn is infeasible. Instead the space lives as a *worklist* of factors $\langle (c_1, W_1), \dots, (c_m, W_m) \rangle$, “draw c_j wires from each W_j ”, and we split recursively [in the manner of maskVerif, Barthe et al. '19]:

- Initially, try to discharge a whole batch at once: if the *union* of its wires is jointly secret-free, every probe inside is also settled — prune the branch.
- On failure, certify one representative *probe* w (one probe of the target size, c_j wires from each factor), grow it into a safe set (Component III), and split every factor along safe / unsafe regions. Recurse on the resulting cells; the all-safe cell needs no visit.

Theorem (exhaustiveness, Ch. 5)

The recursion visits every d -subset exactly once. The proof reads Vandermonde's identity $\binom{m+n}{d} = \sum_k \binom{m}{k} \binom{n}{d-k}$ at the level of *sets*: the cells genuinely partition the probes, they do not merely count them.

Component II: covering all probes

Certifying each d -subset in turn is infeasible. Instead the space lives as a *worklist* of factors $\langle (c_1, W_1), \dots, (c_m, W_m) \rangle$, “draw c_j wires from each W_j ”, and we split recursively [in the manner of maskVerif, Barthe et al. '19]:

- Initially, try to discharge a whole batch at once: if the *union* of its wires is jointly secret-free, every probe inside is also settled — prune the branch.
- On failure, certify one representative *probe* w (one probe of the target size, c_j wires from each factor), grow it into a safe set (Component III), and split every factor along safe / unsafe regions. Recurse on the resulting cells; the all-safe cell needs no visit.


Theorem (exhaustiveness, Ch. 5)

The recursion visits every d -subset exactly once. The proof reads Vandermonde's identity $\binom{m+n}{d} = \sum_k \binom{m}{k} \binom{n}{d-k}$ at the level of *sets*: the cells genuinely partition the probes, they do not merely count them.

The delicate point is that the split must *lose nothing*, and we proved that it does not.

Component III: growing safe sets


A discharged probe is a coupling — but can the *same* coupling cover more wires?



- A wire u joins the safe set of coupling φ when the composite $u \circ \varphi$ is secret-free.

Component III: growing safe sets


A discharged probe is a coupling — but can the *same* coupling cover more wires?



- A wire u joins the safe set of coupling φ when the composite $u \circ \varphi$ is secret-free.
- Exact test: a function depends on a variable iff the variable occurs in its *ANF* — so the endpoint test is to send expressions into their normal forms.

Component III: growing safe sets

A discharged probe is a coupling — but can the *same* coupling cover more wires?



- A wire u joins the safe set of coupling φ when the composite $u \circ \varphi$ is secret-free.
- Exact test: a function depends on a variable iff the variable occurs in its *ANF* — so the endpoint test is to send expressions into their normal forms.
- One certificate, with more wires: the whole safe set is settled by a single discharge, and Component II skips it.

Component III: three mechanisms

- **Closure** — build-time, entropy-style deduction ($H(u \mid w) = 0$). Works locally, but the candidate pool deep in the recursion is scattered. A cheap step but weak.

Proposition (nesting, Ch. 5)

$\text{closure} \subseteq \text{replay} \subseteq \text{coupling}$ — every wire the cheaper mechanisms admit, the coupling admits too. Coppula adopts the coupling approach.

Component III: three mechanisms

- **Closure** — build-time, entropy-style deduction ($H(u \mid w) = 0$). Works locally, but the candidate pool deep in the recursion is scattered. A cheap step but weak.
- **Replay** — re-run the derivation on the candidate wire [after Barthe et al. '15]. An imperfect fit for us because we *factor*, and factoring is not a replayable step.

Proposition (nesting, Ch. 5)

$\text{closure} \subseteq \text{replay} \subseteq \text{coupling}$ — every wire the cheaper mechanisms admit, the coupling admits too. Coppula adopts the coupling approach.

Component III: three mechanisms

- **Closure** — build-time, entropy-style deduction ($H(u \mid w) = 0$). Works locally, but the candidate pool deep in the recursion is scattered. A cheap step but weak.
- **Replay** — re-run the derivation on the candidate wire [after Barthe et al. '15]. An imperfect fit for us because we *factor*, and factoring is not a replayable step.
- **Coupling** — apply the recorded tape bijection and test the ANF endpoint. Exact, and independent of how the derivation was found.

Proposition (nesting, Ch. 5)

$\text{closure} \subseteq \text{replay} \subseteq \text{coupling}$ — every wire the cheaper mechanisms admit, the coupling admits too. Coppula adopts the coupling approach.

Component III: the test

Take the certified probe w with coupling φ , and a candidate wire u from the current pool. Decide whether $u \circ \varphi$ secret-free?

Phase 1 — fast reject (probabilistic, may only *reject*)

Sample tapes; evaluate $u \circ \varphi$ twice per tape, secrets drawn at random vs. all zero. A disagreement proves $u \circ \varphi$ moves with the secrets: drop u . Agreement proves nothing, survivors move on.

Component III: the test

Take the certified probe w with coupling φ , and a candidate wire u from the current pool. Decide whether $u \circ \varphi$ secret-free?

Phase 1 — fast reject (probabilistic, may only *reject*)

Sample tapes; evaluate $u \circ \varphi$ twice per tape, secrets drawn at random vs. all zero. A disagreement proves $u \circ \varphi$ moves with the secrets: drop u . Agreement proves nothing, survivors move on.

Phases 2–3 — exact ($\mathbb{F}_2$ cancellation, then ANF)

A function depends on a variable iff the variable occurs in its unique multilinear normal form. So admit u exactly when no secret survives in the ANF of $u \circ \varphi$ (Lemma, Ch. 5).

Component III: the test

Take the certified probe w with coupling φ , and a candidate wire u from the current pool. Decide whether $u \circ \varphi$ secret-free?

Phase 1 — fast reject (probabilistic, may only *reject*)

Sample tapes; evaluate $u \circ \varphi$ twice per tape, secrets drawn at random vs. all zero. A disagreement proves $u \circ \varphi$ moves with the secrets: drop u . Agreement proves nothing, survivors move on.

Phases 2–3 — exact ($\mathbb{F}_2$ cancellation, then ANF)

A function depends on a variable iff the variable occurs in its unique multilinear normal form. So admit u exactly when no secret survives in the ANF of $u \circ \varphi$ (Lemma, Ch. 5).

Why blinded wires are safe to add. For an admitted u , $u(\varphi_a(\rho); a) = \bar{u}(\rho)$, a function of ρ alone. Since φ_a preserves the uniform law, the *joint* law of all admitted wires under a is that of $(\bar{u}(\rho))_u$, which mentions no secret (Theorem, Ch. 5).

Implementation & Evaluation

Coppula is implemented in **Lean 4**.

- **Hash-consed DAG** expression representation with canonicalization — structural sharing makes equality checks and substitution efficient time- and memory-wise.

Coppula is implemented in **Lean 4**.

- **Hash-consed DAG** expression representation with canonicalization — structural sharing makes equality checks and substitution efficient time- and memory-wise.
- **Multiplicative factoring** as a normal-form pass over the DAG, applied lazily.

Coppula is implemented in **Lean 4**.

- **Hash-consed DAG** expression representation with canonicalization — structural sharing makes equality checks and substitution efficient time- and memory-wise.
- **Multiplicative factoring** as a normal-form pass over the DAG, applied lazily.
- **Bit-sliced ANF endpoint test** for the extension: word-parallel evaluation decides secret-freeness of $u \circ \varphi$ exactly.

Coppula is implemented in **Lean 4**.

- **Hash-consed DAG** expression representation with canonicalization — structural sharing makes equality checks and substitution efficient time- and memory-wise.
- **Multiplicative factoring** as a normal-form pass over the DAG, applied lazily.
- **Bit-sliced ANF endpoint test** for the extension: word-parallel evaluation decides secret-freeness of $u \circ \varphi$ exactly.

Design and proofs are implementation-agnostic; these choices simply make algorithms fast.

Evaluation: outcomes

Forty configurations at orders 1–4: ISW multiplications, DOM and TI ANDs, refresh gadgets, Keccak χ (DOM and TI sharings), and the TI S-boxes and quadratic sharings of Bilgin et al. Every verdict matches the ground truth, and every Insecure comes with a concrete witness observation.

- **DOM-AND** (2 shares): Secure at order 1, the witness pair $\{c_{01}, c_{10}\}$ at order 2 — our running example.

Evaluation: outcomes

Forty configurations at orders 1–4: ISW multiplications, DOM and TI ANDs, refresh gadgets, Keccak χ (DOM and TI sharings), and the TI S-boxes and quadratic sharings of Bilgin et al. Every verdict matches the ground truth, and every Insecure comes with a concrete witness observation.

- **DOM-AND** (2 shares): Secure at order 1, the witness pair $\{c_{01}, c_{10}\}$ at order 2 — our running example.
- **Four-share $x \oplus yz$ TI** [Bilgin et al. '15]: certified secure at orders 1 and 2, insecure at order 3 — the *expected* behaviour of a first-order threshold implementation.

Evaluation: outcomes

Forty configurations at orders 1–4: ISW multiplications, DOM and TI ANDs, refresh gadgets, Keccak χ (DOM and TI sharings), and the TI S-boxes and quadratic sharings of Bilgin et al. Every verdict matches the ground truth, and every Insecure comes with a concrete witness observation.

- **DOM-AND** (2 shares): Secure at order 1, the witness pair $\{c_{01}, c_{10}\}$ at order 2 — our running example.
- **Four-share $x \oplus yz$ TI** [Bilgin et al. '15]: certified secure at orders 1 and 2, insecure at order 3 — the *expected* behaviour of a first-order threshold implementation.
- **The factoring family** ($d+1$ quadratic S-boxes, Q^3/Q^4): randomness-free sharings whose certification must factor first. A factoring-free, maskVerif-style replay engine false-reports Insecure on five of them; Coppula verifies every one — and the two it does report insecure are exactly the quadratic classes with no uniform three-share TI.

Evaluation: outcomes


Forty configurations at orders 1–4: ISW multiplications, DOM and TI ANDs, refresh gadgets, Keccak χ (DOM and TI sharings), and the TI S-boxes and quadratic sharings of Bilgin et al. Every verdict matches the ground truth, and every Insecure comes with a concrete witness observation.

- **DOM-AND** (2 shares): Secure at order 1, the witness pair $\{c_{01}, c_{10}\}$ at order 2 — our running example.
- **Four-share $x \oplus yz$ TI** [Bilgin et al. '15]: certified secure at orders 1 and 2, insecure at order 3 — the *expected* behaviour of a first-order threshold implementation.
- **The factoring family** ($d+1$ quadratic S-boxes, Q^3/Q^4): randomness-free sharings whose certification must factor first. A factoring-free, maskVerif-style replay engine false-reports Insecure on five of them; Coppula verifies every one — and the two it does report insecure are exactly the quadratic classes with no uniform three-share TI.

(Full extension-comparison table @ the appendix slides.)


Evaluation: is the extension worth it?

- Coupling matches replay's yield *to the digit* wherever replay is right (five-share ISW: 4735 discharges, 8078 free wires on both) — and keeps the correct verdict on the five rows where replay false-reports Insecure.



Verification time versus circuit depth.


Evaluation: is the extension worth it?



Verification time versus circuit depth.

- Coupling matches replay's yield *to the digit* wherever replay is right (five-share ISW: 4735 discharges, 8078 free wires on both) — and keeps the correct verdict on the five rows where replay false-reports Insecure.
- Closures keep the verdicts but lose the yield: 14 vs. 537 free wires on Fides, 632,834 vs. 4735 discharges on five-share ISW — the split degrades toward enumeration.

Evaluation: is the extension worth it?



Verification time versus circuit depth.

- Coupling matches replay's yield *to the digit* wherever replay is right (five-share ISW: 4735 discharges, 8078 free wires on both) — and keeps the correct verdict on the five rows where replay false-reports Insecure.
- Closures keep the verdicts but lose the yield: 14 vs. 537 free wires on Fides, 632,834 vs. 4735 discharges on five-share ISW — the split degrades toward enumeration.
- The extension pays for itself: Fides certifies 537 wires from 55 discharges; five-share ISW falls from ~ 609 s (no extension) to under 2 s.

Wrap-up

- **Model.** Our model is value-based leakage only: glitches and transitions are out of scope; so is composition. So we show single-circuit, exact probing security.

Limitations and future work

- **Model.** Our model is value-based leakage only: glitches and transitions are out of scope; so is composition. So we show single-circuit, exact probing security.
- **Larger fields.** The core — shears, couplings, ANF — is already field-agnostic; lifting the implementation beyond $\mathbb{F}_2$ is engineering plus a normal-form choice. (Also, the power of Schwartz–Zippel can come into picture here.)

Limitations and future work

- **Model.** Our model is value-based leakage only: glitches and transitions are out of scope; so is composition. So we show single-circuit, exact probing security.
- **Larger fields.** The core — shears, couplings, ANF — is already field-agnostic; lifting the implementation beyond $\mathbb{F}_2$ is engineering plus a normal-form choice. (Also, the power of Schwartz–Zippel can come into picture here.)
- **Composing couplings.** for unions of safe sets using couplings as a guide.

Limitations and future work

- **Model.** Our model is value-based leakage only: glitches and transitions are out of scope; so is composition. So we show single-circuit, exact probing security.
- **Larger fields.** The core — shears, couplings, ANF — is already field-agnostic; lifting the implementation beyond $\mathbb{F}_2$ is engineering plus a normal-form choice. (Also, the power of Schwartz–Zippel can come into picture here.)
- **Composing couplings.** for unions of safe sets using couplings as a guide.
- **More.** smarter factoring heuristics; exploiting circuit symmetries to quotient the subset space.

1. **Placement.** We show that: counting is $\#P$ -complete, deciding is $C=P$ -complete — above the polynomial hierarchy, ruling out the counting route.

1. **Placement.** We show that: counting is $\#P$ -complete, deciding is $C=P$ -complete — above the polynomial hierarchy, ruling out the counting route.
2. **A rigorous coupling account of probing security.** Security *is* the existence of a family of coupling functions, the family is a groupoid, and substitutions construct it.

1. **Placement.** We show that: counting is $\#P$ -complete, deciding is $C=P$ -complete — above the polynomial hierarchy, ruling out the counting route.
2. **A rigorous coupling account of probing security.** Security *is* the existence of a family of coupling functions, the family is a groupoid, and substitutions construct it.
3. **Coppula.** Made up of a sound rewrite engine (with factoring), an exhaustive subset-space traversal, and a probe-set extension with a proven mechanism hierarchy closure $\subseteq$ replay $\subseteq$ coupling.

1. **Placement.** We show that: counting is $\#P$ -complete, deciding is $C=P$ -complete — above the polynomial hierarchy, ruling out the counting route.
2. **A rigorous coupling account of probing security.** Security *is* the existence of a family of coupling functions, the family is a groupoid, and substitutions construct it.
3. **Coppula.** Made up of a sound rewrite engine (with factoring), an exhaustive subset-space traversal, and a probe-set extension with a proven mechanism hierarchy closure $\subseteq$ replay $\subseteq$ coupling.
4. **Evaluation.** We present exact verdicts, with witnesses, across a suite of masked gadgets, and the coupling extension wins in coverage.

Thank you!

Questions?

Backup

Backup: extension comparison (I/IV)

Each entry: Coupling / Replay / Closure. Verdict S = Secure, I = Insecure; † marks replay's false-Insecures.

Circuit	d	Verdict	Time [ms]	Discharges	Free wires
ISW multiplication					
2 sh.	1	S/S/S	0/0/1	7/7/15	2/2/0
3 sh.	2	S/S/S	3/1/34	34/34/309	31/31/8
4 sh.	3	S/S/S	46/37/3190	338/338/10770	437/437/439
5 sh.	4	S/S/S	1125/845/409472	4735/4735/632834	8078/8078/37513
DOM AND					
2 sh.	1	S/S/S	1/0/0	11/11/15	2/2/0
2 sh.	2	I/I/I	0/0/1	2/2/6	0/0/2
3 sh.	2	S/S/S	5/4/35	88/88/325	44/44/7
4 sh.	3	S/S/S	104/90/3419	859/859/13291	784/784/532

Backup: extension comparison (II/IV)

Circuit	d	Verdict	Time [ms]	Discharges	Free wires
Trichina AND					
well-formed	1	S/S/S	0/1/1	11/11/15	2/2/0
mis-associated	1	I/I/I	1/0/1	10/10/12	1/1/0
TI AND					
3 sh.	1	S/S/S	0/0/1	9/9/27	9/9/0
3 sh.	2	I/I/I	1/1/1	10/10/10	5/5/2
Mask refreshing					
additive, 3 sh.	2	S/S/S	0/0/0	6/6/8	1/1/0
additive, 4 sh.	3	S/S/S	1/0/1	11/11/17	6/6/0
ISW, 3 sh.	2	S/S/S	0/0/0	6/6/10	4/4/0
ISW, 4 sh.	3	S/S/S	0/1/2	12/12/51	24/24/0
Keccak χ					
DOM, 2 sh.	1	S/S/S	14/11/29	49/49/107	29/29/5
DOM, 3 sh.	2	S/S/S	2807/1797/11563	2865/2865/14047	2671/2671/645
TI, 3 sh.	1	S/S/S	21/14/103	43/43/183	68/68/5
TI, 3 sh.	2	I/I/I	21/14/72	57/57/352	92/92/27

Backup: extension comparison (III/IV)

Circuit	d	Verdict	Time [ms]	Discharges	Free wires
TI S-boxes and quadratic sharings					
Q_{12}^4 direct, 2 sh.	1	S/I [†] /S	2/2/5	27/26/47	16/16/10
Q_{12}^4 TI, 3 sh.	1	S/S/S	5/4/25	25/25/83	30/30/6
Fides AB1, 4 sh.	1	S/S/S	1941/972/36461	55/55/1127	537/537/14
$x + yz$ direct, 3 sh.	1	S/S/S	1/0/3	7/7/29	10/10/0
$x + yz$ w/ CT, 3 sh.	1	S/S/S	1/1/3	11/11/29	9/9/0
$x + yz$, 4 sh.	2	S/S/S	11/8/170	78/78/767	115/115/20
$x + yz$, 4 sh.	3	I/I/I	11/7/8	102/102/72	99/99/15
$x + yz + xyz$, 4 sh.	1	S/S/S	45/18/175	31/31/167	63/63/1
xy virtual var.	1	S/S/S	3/1/17	9/9/55	23/23/0
xy virtual share	1	S/S/S	1/0/7	9/9/37	15/15/0
xy , $4 \rightarrow 3$ sh.	1	S/S/S	1/1/3	11/11/25	8/8/0

Backup: extension comparison (IV/IV)

Circuit	d	Verdict	Time [ms]	Discharges	Free wires
$d+1$ quadratic S-box family					
Q_1^3	1	S/S/S	0/1/0	13/13/23	11/11/10
Q_2^3	1	S/I ⁺ /S	2/1/4	25/24/43	13/13/6
Q_3^3	1	I/I/I	4/2/7	18/18/32	22/22/2
Q_4^4	1	S/S/S	1/0/1	15/15/27	14/14/14
Q_{12}^4	1	S/I ⁺ /S	2/2/5	25/24/47	15/15/8
Q_{293}^4	1	S/I ⁺ /S	4/2/6	31/24/59	20/20/8
Q_{294}^4	1	S/S/S	1/1/3	23/23/39	16/16/12
Q_{299}^4	1	S/I ⁺ /S	5/2/13	35/18/75	26/22/8
Q_{300}^4	1	I/I/I	4/3/8	18/18/32	24/24/2

The two genuinely insecure classes, Q_3^3 and Q_{300}^4 , are exactly the quadratic classes that admit no uniform three-share TI [Bilgin et al. '15].

Backup: C=P membership, sketch

Encode disagreement of the two laws as a single gap function:

$$f(C, O, a, b) = \sum_o (N_a(o) - N_b(o))^2.$$

Each N is a #P function; GapP is closed under subtraction and multiplication [Fenner et al. '94]; the sum has $2^k \leq 2^d$ terms, constant for fixed d . A sum of squares vanishes iff every term does, so

$$\mathcal{L}_a(O) = \mathcal{L}_b(O) \iff f = 0,$$

and the problem is the vanishing set of a gap function: C=P.

Backup: replay vs. factoring

Barthe et al.'s derivations replay two rule kinds: OPT (substitution) and CONV (ANF normalisation) — for them, ANF is a *replayed rule*, not a fallback.

Our derivations interleave substitutions with *multiplicative factoring*, and a factoring step has no tape-level counterpart to replay. Replaying the substitution steps alone on a new wire can spuriously fail, producing false-Insecure verdicts (observed in the evaluation, the † rows).

Hence the coupling mechanism: apply the recorded bijection $\widehat{\sigma}_1 \circ \dots \circ \widehat{\sigma}_k$ semantically and test the ANF endpoint — exact, and indifferent to how the derivation was found.

Backup: positioning

- **maskVerif** [Barthe et al.]: language-based, the traversal's main inspiration; replay-style extension; no factoring.
- **REBECCA** / **Fourier-analytic** [Bloem et al.]: correlation sets over Fourier coefficients; approximate in general.
- **SILVER** [Knichel et al.]: exact, BDD-based model counting; exhaustive per-probe distributions, different scaling profile.
- Coppula: exact like SILVER, constructive like maskVerif — and the certificate (the coupling) is itself an object with structure we exploit.

Backup: the four-share TI netlist

Four-share TI of $x \oplus (y \odot z)$ [Bilgin et al. '15]; each input four-shared, sixteen cross products $y_i \odot z_j$:

$$\begin{aligned} F_1 &= x_2 \oplus \bigoplus_{i,j \in \{2,3,4\}} y_i z_j, & F_2 &= x_3 \oplus y_1 z_3 \oplus y_1 z_4 \oplus y_3 z_1 \oplus y_4 z_1 \oplus y_1 z_1, \\ F_3 &= x_4 \oplus y_1 z_2 \oplus y_2 z_1, & F_4 &= x_1. \end{aligned}$$

A *first-order* threshold implementation — never claimed beyond order one. Coppula certifies orders 1 and 2 and finds the order-3 witness: a counterexample on a gadget with a *known* failure order, not a flaw.